

Lecture 03: Contiguous Tables + Aggregation Pipelines

Goals

- Model data as tables (rows + columns)
- Build extraction -> normalization -> aggregation pipelines
- Avoid misaligned columns in parallel arrays

Techniques / new functions

- `Path`, `csv.DictReader`, list/dict comprehensions
- `str.strip()`, `str.casefold()`, defaulting with `or`
- `list.index()`, `zip()`, `len()`, `all()`
- `collections.Counter.most_common()`

Data recap (Spotify)

Demo files live in `lectures/data/` (or `data/` if you open this notebook inside `lectures/`). We will reformat list-of-dicts into table form for faster scans and summaries.

Segment 1: Table mindset

Goal: rows + columns with a header.

Techniques / new functions

- Build tables with `header` + row lists
- Read rows with `csv.DictReader`
- Normalize text via `strip()` and `casefold()`

In []:

```
from pathlib import Path
import csv

data_dir = Path('lectures/data')
if not data_dir.exists():
    data_dir = Path('data')

def parse_int(value, default=0):
    try:
        return int(value)
    except (TypeError, ValueError):
        return default

def clean_track_row(row):
    return {
        'track_id': row['track_id'].strip().casefold(),
        'track_name': row['track_name'].strip(),
        'artist_id': row['artist_id'].strip().casefold(),
        'genre': (row['genre'] or 'unknown').strip().casefold(),
        'duration_ms': parse_int(row['duration_ms']),
    }
```

```
In [ ]: with (data_dir / 'tracks.csv').open(newline='', encoding='utf-8') as f:
        reader = csv.DictReader(f)
        tracks = [clean_track_row(r) for r in reader]

        print(tracks)
        header = ['track_id', 'track_name', 'artist_id', 'genre', 'duration_ms']
        table = [header] + [[t[col] for col in header] for t in tracks]
        table[:4]
```

Segment 2: Parallel arrays

Goal: keep aligned lists for fast scans.

Techniques / new functions

- Extract columns using `list.index()`
- Keep arrays aligned by index
- Recombine with `zip()` for inspection

```
In [ ]: def column(table, name):  
        idx = table[0].index(name)  
        return [row[idx] for row in table[1:]]  
  
column(table, 'genre')
```

Spotlight: `zip()`

Definition: `zip()` iterates multiple sequences in lockstep, producing tuples of aligned items.

Usage: `zip(a, b, c)` yields `(a[i], b[i], c[i])` until the shortest input ends.

Common reasons / use cases

- Recombine parallel arrays to inspect or iterate aligned rows
- Build rows for tables from separate columns
- Pair IDs with values (e.g., keys with counts) without manual indexing

```
In [ ]: track_ids = column(table, 'track_id')
genres = column(table, 'genre')
durations = column(table, 'duration_ms')

print(track_ids)
print(genres)
print(durations)
list(zip(track_ids, genres, durations))
```


Alignment check

- All arrays must have the same length
- The same index points to the same track

Segment 3: Pipeline stages

Goal: extract -> normalize -> aggregate -> report.

Techniques / new functions

- Write stage functions (`extract_*`, `normalize_*`, `report_*`)
- Aggregate with `collections.Counter`
- Return top-k using `most_common()`

```
In [ ]: def extract_tracks(rows):  
        return [[r['track_id'], r['artist_id'], r['genre'], r['duration_ms']]  
  
        def normalize_rows(rows):  
            return [[tid.strip().casefold(), aid.strip().casefold(), (g or 'unknown'), dur]   
                    for tid, aid, g, dur in rows]  
  
        def report_top_genres(rows, k=3):  
            from collections import Counter  
            counts = Counter(r[2] for r in rows)  
            return counts.most_common(k)  
  
        stage1 = extract_tracks(tracks)  
        stage2 = normalize_rows(stage1)  
        print(stage1)  
        print(stage2)  
        report_top_genres(stage2)
```

Segment 1-3 summary

- Tables make scans predictable
- Parallel arrays are fast but easy to misalign
- Pipelines keep each step focused

In-class exercise 1 (commit)

Build a `table` with a header row for plays (`play_id`, `track_id`, `play_count`).

- Write a `column_sum(table, 'play_count')` helper
- Commit your solution

Break (3 minutes)

Segment 4: Artist-by-genre table

Goal: produce a 2D summary table.

Techniques / new functions

- Iterate `dict.items()` to build rows
- Add a header row to a table
- Append rows with `.append()`

```
In [ ]: def clean_play_row(row):  
        return [  
            row['play_id'].strip().casefold(),  
            row['track_id'].strip().casefold(),  
            parse_int(row['play_count']),  
        ]  
  
        with (data_dir / 'plays.csv').open(newline='', encoding='utf-8') as f:  
            reader = csv.DictReader(f)  
            plays = [clean_play_row(r) for r in reader]  
  
        plays[:4]
```



```
In [ ]: # Convert summary dict to a table with header
summary_table = [['artist_id', 'genre', 'total_plays']]
for (artist_id, genre), total in summary.items():
    summary_table.append([artist_id, genre, total])

summary_table
```

Segment 5: Verify table shape

Goal: catch misaligned columns early.

Techniques / new functions

- Validate width with `len()`
- Check all rows using `all()`

Spotlight: `all()`

Definition: `all()` returns `True` if every item in an iterable is truthy.

Usage: `all(condition for x in items)` checks a condition across the whole sequence.

Common reasons / use cases

- Validate that every row in a table has the same length
- Assert data quality checks across a dataset (e.g., non-empty fields)
- Combine multiple boolean checks into one guard

```
In [ ]: def verify_table(table):  
        width = len(table[0])  
        return all(len(row) == width for row in table)  
  
        verify_table(summary_table)
```

Segment 6: Tradeoffs

Goal: know when tables beat dicts.

Techniques / new functions

- Choose tables for fast scans + predictable order
- Choose dicts for keyed lookups + sparse data

- Tables are great for scans, sorting, and exporting
- Dicts are great for lookups and sparse data
- Choose the shape that matches your dominant operation

Segment 4-6 summary

- Build summary tables from dict aggregations
- Verify shapes to avoid silent errors
- Tables and dicts complement each other

In-class exercise 2 (commit)

Create a 2D table for `artist_id`, `track_count`, `avg_duration_ms`.

- Use the tracks list
- Verify the table shape
- Commit your results

Wrap-up

- Tables + pipelines make analytics predictable
- Keep columns aligned and verify shapes
- Next time: hash indexing for fast joins

